

CASA: A Context-Aware Security Pipeline for Personal AI Agents

1st Liu, Chia-Chiao

Dept. of Computer Science and Engineering

Texas A&M University

College Station, The United States

ccl@tamu.edu

UIN: 637004273

Abstract—AI agents powered by large language models are increasingly deployed as autonomous software entities capable of interacting with operating systems through tool use and system-level actions. While this capability enables powerful automation, it also introduces new security risks, as high-level instructions are dynamically translated into multi-step execution sequences that are difficult to monitor and control using existing security mechanisms.

In this work, we present CASA, a context-aware security architecture for AI agents that bridges low-level system observability with high-level behavioral reasoning. CASA leverages eBPF to efficiently capture system events such as process execution, file access, and network activity, and reconstructs structured execution context in user space, including process lineage, privilege state, and behavioral history. A programmable rule engine then evaluates this context to produce risk scores over multi-step behaviors, enabling flexible and expressive security policies.

Unlike existing approaches that operate on isolated system calls or static attributes, CASA reasons over session-level execution context, allowing it to capture cross-event and context-dependent behaviors. The system is non-invasive and agent-agnostic, requiring no modification to AI agents and operating entirely at the operating system boundary.

We demonstrate that context-aware reconstruction of execution behavior provides a practical foundation for securing AI agent workloads in dynamic and open-ended environments.

Index Terms—AI Agents, System Security, eBPF, Runtime Monitoring, Context-Aware Security, Behavioral Analysis, Programmable Policy

I. INTRODUCTION

The emergence of Large Language Model (LLM)-based agents represents a fundamental shift in how software interacts with operating systems. Unlike traditional software, where system-level actions are explicitly triggered by human users, AI agents autonomously generate and execute actions through language-driven reasoning. In conventional settings, a human user invokes commands through applications or shells, resulting in system calls with relatively clear intent and predictable behavior. Security policies in such environments can therefore be defined a priori, based on known workflows and expected usage patterns.

In contrast, AI agents operate under a non-deterministic execution model where high-level intent is decomposed into a sequence of runtime-generated actions. Unlike atomic system calls, these actions are inherently multi-step, context-dependent, and often unpredictable at design time. A single




Fig. 1. Traditional system-level monitoring observes isolated system calls, while context-aware analysis reconstructs higher-level behavioral semantics.

instruction may translate into a chain of system interactions whose semantics only emerge at runtime. As a result, traditional security mechanisms that operate on individual system calls or rely on static policy definitions struggle to capture the full intent behind an agent's behavior.

This challenge is particularly pronounced in the setting of personal AI agents. While enterprise systems often rely on predefined workflows and strict access control boundaries, personal AI agents operate in open-ended environments with highly dynamic and user-driven tasks. They are commonly deployed with broad permissions over local resources, including file systems, network interfaces, and system utilities. However, this flexibility makes it difficult to define fixed security policies in advance, as legitimate behavior may vary significantly across tasks and over time. Consequently, security decisions must be adaptive and informed by execution context rather than static rules alone.

At the system level, existing security mechanisms provide only a fragmented view of execution. Modern observability tools offer fine-grained visibility into system calls and process activities, while enforcement mechanisms such as sandboxing and system call filtering constrain behavior based on predefined attributes. However, these approaches fundamentally operate on isolated events, treating each system call as an independent unit of analysis.

This fragmentation leads to a semantic gap between low-level observations and high-level behavior. As illustrated in Figure 1, sequences of system calls such as connect, openat, write, and execve are observed as disjoint events, without




Fig. 2. Multi-stage breakdown of a prompt injection attack. While the agent follows a coherent logical path (left), the underlying operating system perceives only a series of fragmented, low-level system calls (right). Without cross-event correlation, existing security tools fail to distinguish these malicious steps from legitimate tool-use behaviors.

capturing their semantic relationships. However, AI agent behavior emerges from the composition of such events over time, forming meaningful multi-step actions such as download-and-execute. Bridging this gap requires reasoning beyond individual events toward structured, context-aware interpretations of execution.

This limitation becomes particularly evident in realistic attack scenarios involving AI agents. Consider a prompt injection attack in which a user requests the agent to summarize a webpage, but the content includes hidden instructions to download and execute a remote payload. As shown in Figure 2, this results in a sequence of system-level actions including network connection (connect), file retrieval (open/write), and execution (execve). While each individual action is benign in isolation, their composition forms a clear malicious pattern that cannot be captured by event-level detection or static rules.

Motivated by this gap, we design a context-aware security pipeline that combines low-level event capture with user-space context reconstruction and programmable policy evaluation, enabling reasoning over multi-step AI agent behaviors.

To address this gap, we propose a context-aware security pipeline that reconstructs execution behavior from low-level system events and enables reasoning over multi-step agent actions.

We instantiate this idea in CASA, a context-aware security architecture for AI agents that integrates kernel-level observability with user-space context reconstruction and programmable policy evaluation. Rather than inspecting isolated system calls, CASA organizes analysis as a two-plane pipeline that connects runtime telemetry collection with session tracking, structured context derivation, and risk-based decision making.

In the data plane, CASA combines eBPF-based event capture with process-state inspection to obtain both dynamic execution signals and static security attributes. In the control plane, CASA groups events into execution sessions, derives structured context along execution, capability, and history

dimensions, and evaluates this context using a programmable rule engine to produce risk scores over multi-step behaviors.

This design enables CASA to capture AI-agent-specific attack patterns, such as prompt injection leading to tool misuse and command execution, while remaining lightweight, modular, and applicable to open-ended runtime environments.

Importantly, our approach is non-invasive and agent-agnostic, operating entirely at the operating system boundary without requiring any modification to the AI agent, and enabling practical deployment across diverse agent environments.

We argue that securing AI agents fundamentally requires moving beyond event-level detection toward context-aware behavioral reasoning, while remaining programmable and practical for real-world use.

We summarize our contributions as follows:

- **Context-aware behavioral analysis pipeline.** We design a system that reconstructs execution context across events, processes, and time, enabling reasoning over multi-step behaviors that are difficult to express using traditional event-level security mechanisms.
- **Programmable policy framework over structured context.** We provide a flexible rule engine that allows security policies to be defined over derived execution context, supporting expressive and composable detection logic with minimal engineering overhead.
- **System-level modeling of AI-agent attack patterns.** We identify and characterize representative attack scenarios in AI-agent environments, such as prompt injection leading to tool misuse and command execution, and show how these behaviors manifest as multi-step system-level activity.
- **Non-invasive and agent-agnostic deployment.** Our approach operates entirely at the operating system boundary using eBPF-based observability, requiring no modification to the AI agent and enabling applicability across diverse agent frameworks.

II. RELATED WORK

A. Runtime Security Systems

Modern runtime security systems leverage kernel-level observability to monitor system behavior and detect suspicious activities. Tools such as Falco [1] and Tetragon [2] utilize eBPF to capture system calls, process events, and network activities in real time. These systems typically apply rule-based detection over observed events, allowing users to specify conditions under which alerts should be triggered, such as unexpected process execution or access to sensitive files.

The primary strength of these systems lies in their ability to provide fine-grained visibility into low-level system activity with relatively low overhead. By operating at the kernel level, they can observe all processes uniformly without requiring instrumentation of individual applications.

As a result, these systems struggle to capture relationships across events and processes, making it difficult to express

multi-step behaviors or reconstruct higher-level execution semantics.

B. Kernel Enforcement Mechanisms

Operating systems also provide built-in enforcement mechanisms to restrict process behavior at runtime. Notable examples include `seccomp` [3], which filters system calls based on predefined policies, and the Linux Security Module (LSM) framework [4], which enables fine-grained access control through pluggable security modules.

These mechanisms offer strong security guarantees due to their position within the kernel. They enable immediate enforcement decisions at the point of execution, ensuring that disallowed actions are blocked before they are carried out, while maintaining low runtime overhead.

Despite these advantages, kernel enforcement mechanisms are inherently constrained in expressiveness. Their policies are primarily attribute-based, relying on static properties such as system call type or process identity. They do not maintain historical execution state or capture behavioral context across multiple events. Consequently, they are unable to express policies that depend on multi-step behaviors or evolving execution semantics.

Consequently, they cannot express policies that depend on execution history or cross-event relationships, limiting their effectiveness in detecting complex, multi-step behaviors.

C. Programmable Observability and Policy Systems

Recent systems research has explored programmable observability, where users define custom logic over runtime telemetry. Frameworks such as `OpenTelemetry` [5] and `eBPF`-based tracing pipelines enable flexible processing of system signals for monitoring and debugging purposes.

The key advantage of this approach is flexibility. Unlike fixed rule-based systems, programmable pipelines can express more complex logic and adapt to changing workloads.

However, these systems lack security-oriented abstractions and do not model adversarial behavior, making them insufficient for detecting structured attack patterns in AI-agent workloads.

D. AI Agent Security

Security concerns surrounding AI agents have recently gained attention, particularly in the context of prompt injection attacks and tool misuse [6]. These attacks exploit the agent’s reliance on external inputs and tool invocation mechanisms to induce unintended behaviors.

Existing defenses are largely implemented at the application or model layer, focusing on input filtering, prompt sanitization, and alignment-based techniques. While effective in limiting direct exploitation, these approaches assume correct model behavior and do not account for system-level execution effects.

In contrast, there is limited work addressing system-level monitoring or enforcement for AI agents. This leaves a critical gap, as many attacks ultimately manifest as sequences of system-level actions that are invisible to application-layer defenses.

As a result, system-level execution effects remain largely unmonitored, leaving a gap between application-layer defenses and underlying system behavior.

E. The Semantic Gap in System-Level Security

Despite extensive work in runtime security, kernel enforcement, and observability systems, existing approaches are not designed for the execution model introduced by AI agents. In particular, AI agents transform high-level language instructions into multi-step, tool-mediated system interactions that unfold dynamically at runtime. This introduces a fundamentally different threat surface where malicious behavior is not expressed as a single event, but as a sequence of individually benign operations whose semantics only emerge at the behavioral level.

As a result, prior systems that operate on isolated system calls, static attributes, or short event patterns are unable to reconstruct the full execution context required to reason about agent behavior. This mismatch between event-level observability and behavior-level semantics highlights a fundamental limitation of existing systems: they lack the ability to reconstruct execution semantics across events, processes, and time.

F. Summary

In contrast to prior work, our approach explicitly bridges kernel-level observability and behavior-level reasoning by reconstructing structured execution context and enabling programmable policies over multi-step behaviors. This allows us to capture execution patterns that are fundamentally difficult to express in existing event-centric or attribute-based systems.

III. SYSTEM DESIGN

CASA is designed around three goals. First, it should reconstruct multi-step execution behavior rather than analyze isolated events. Second, it should remain non-invasive and agent-agnostic by operating at the operating system boundary. Third, it should support programmable policies so that security logic can be adapted to different agent environments without modifying the core system.

A. System Overview

We design CASA, a context-aware security architecture that analyzes the runtime behavior of AI agents by transforming low-level system events into structured, session-level execution context, and evaluating this context using programmable security rules.

Our current prototype targets the **OpenClaw runtime** as a representative personal-agent environment. In particular, the implementation monitors the `openclaw-gateway` process tree and relies on runtime-specific session boundaries to associate low-level activity with individual agent-facing **CLI invocations**.

Fig. 3 illustrates the overall system architecture. At a high level, the system is organized into two planes: a **Data Plane** responsible for collecting raw execution signals, and a **Control Plane** responsible for tracking sessions, reconstructing context, and performing security analysis.




Fig. 3. CASA two-plane analysis pipeline. The data plane collects eBPF Events and process-state information through a /proc Reader, while the control plane tracks sessions, derives structured context, evaluates rules, and emits audit decisions.

The Data Plane consists of two complementary sources. First, **eBPF Events** provide a continuous stream of system-level activity, including process execution, file access, network communication, and process exit. Second, a **/proc Reader** retrieves process-state information from the operating system, such as capability sets and seccomp mode, supplementing the event stream with security-relevant attributes that are not directly observable from system calls alone. Together, these sources provide both dynamic and static views of process behavior.

The Control Plane processes these inputs through a multi-stage pipeline. Incoming events are first ingested by the **Event Tracker** and then grouped by the **Session Tracker** into coherent execution sessions that span related processes and execution steps. The resulting session state is maintained in the **Context Store**, from which the **Context Builder** derives structured context along three dimensions: **execution**, **capability**, and **history**. The **Rule Engine** evaluates this derived context using a configurable risk model, producing a numerical risk score for each session. Finally, the **Decision Engine** maps this score to actions such as **Logging/Alerting**.

This design separates event collection, context reconstruction, and policy evaluation into modular components, enabling expressive and flexible reasoning over complex behaviors while maintaining a lightweight and non-invasive integration with the underlying system.

B. Threat Model

We consider threat scenarios in which an AI agent is partially compromised through indirect means, such as prompt injection, malicious tool outputs, or unintended instruction execution. In these cases, the agent continues to operate within its normal environment but may perform actions that deviate from its intended behavior.

The attacker does not require direct system-level access. Instead, they influence the agent through inputs that cause it to generate and execute unintended commands. As a result, the

agent may invoke system-level operations including process execution, file access, and network communication.

A defining characteristic of such attacks is that they unfold as multi-step sequences of actions, rather than single anomalous events. For example, an injected instruction may cause an agent to download a remote payload and subsequently execute it. Each individual operation, such as establishing a network connection or executing a binary, may appear benign in isolation. However, when considered as part of a sequence, they reveal a clear malicious pattern.

We assume that the underlying operating system and kernel are trusted, and that the attacker does not exploit kernel vulnerabilities or gain direct root-level control outside the agent’s execution context. Our goal is therefore not to prevent all forms of compromise, but to detect and characterize suspicious behaviors at runtime, based on how actions are composed over time and across processes.

This threat model motivates a shift from event-level detection to context-aware analysis, where security decisions are based on aggregated behavior rather than individual system calls.

C. Session Model

A central design choice in our system is the introduction of a session model as the primary unit of analysis. Rather than evaluating system events in isolation, we group related events into sessions that represent coherent units of execution.

A session is defined as a collection of system events that are logically connected through process relationships and execution flow. Specifically, events are grouped based on process identity and extended through process lineage, including execve-based image replacement and related child-process activity. When a process spawns a child or replaces itself with a new executable, the resulting activity is treated as a continuation of the same session. This allows the system to track behavior across multiple execution steps, even when they involve different binaries or subprocesses.

The session model serves two purposes. First, it enables aggregation of events across time, allowing the system to observe sequences of actions rather than isolated operations. Second, it provides a mechanism to preserve execution continuity, ensuring that behaviors are interpreted within the context of how they were initiated and evolved.

This abstraction is particularly important in AI agent environments, where a single high-level task may result in a chain of system-level operations spanning multiple processes and execution stages. By organizing events into sessions, the system captures these behaviors as unified entities, enabling more accurate and meaningful analysis.

D. Context Representation

To support context-aware analysis, the system represents each session using a structured set of features that capture different aspects of execution. This derived context serves as the input to the rule engine and forms the basis for security decisions.

Rather than relying on raw events, the system derives higher-level context along three complementary dimensions: execution, capability, and history. Each dimension captures a distinct aspect of system activity, and together they provide a comprehensive view of how a session evolves. In the implementation, these dimensions are materialized as structured feature fields such as `execution.deep_chain`, `capability.seccomp_disabled`, and `history.connect_then_exec`.

This design reflects the observation that meaningful security signals in AI-agent environments arise not from individual actions, but from the interaction between origin, privilege, and behavior over time. By explicitly modeling these dimensions, the system enables structured reasoning about complex execution patterns.

1) *Execution Chain Context*: Execution chain context captures the origin and lineage of processes within a session. It represents the sequence of process relationships that lead to a given execution, including parent-child links established through process creation and replacement.

This context allows the system to reason about how an action was initiated. In many cases, the legitimacy of an operation depends not only on the action itself, but also on the path through which it was reached. For example, a process launched directly by a trusted service may be expected, while one originating from a transient or untrusted execution path may indicate anomalous behavior.

By modeling execution chains, the system can detect irregular patterns such as unusually deep process hierarchies, unexpected parent processes, or execution paths involving temporary or in-memory artifacts. In the current prototype, representative execution features include suspicious-path execution, deep process lineage, shell presence in the lineage, network-tool presence, interpreter presence, container-runtime presence, and execution from deleted or `memfd`-backed binaries. These signals are particularly relevant in identifying indirect or staged attacks, where malicious behavior is introduced through intermediate steps.

2) *Capability Context*: Capability context represents the privilege state of processes within a session, focusing on the effective capabilities available at runtime. These capabilities determine what operations a process is allowed to perform, including actions that bypass standard access control mechanisms.

This context enables reasoning about the potential impact of behavior. The same action may have different security implications depending on the privileges of the process performing it. For instance, a network operation performed by a highly privileged process may carry greater risk than one performed by a restricted process.

In addition to capabilities, the system also considers related security attributes such as `seccomp` mode, which indicates the degree of `syscall` restriction applied to a process. The current prototype derives capability-oriented features such as whether dangerous capabilities are present, how many dangerous capabilities are active, and whether `seccomp` is disabled.

By incorporating these signals, the system captures a more complete view of the execution environment and its associated risk.

3) *Historical Context*: Historical context captures the temporal evolution of a session, recording the sequence of actions performed over time. This includes system events such as file access, network communication, and process execution, as well as derived patterns that emerge from their combination.

This context enables reasoning about behavioral patterns, particularly those that are only meaningful when viewed across multiple steps. Many attacks in AI-agent environments are inherently sequential, where the combination of actions, rather than any single action, indicates malicious intent. For example, a network connection followed by execution of a newly retrieved binary may signal a potential compromise.

By maintaining historical context, the system can identify such patterns and distinguish between benign and suspicious sequences of activity. In the current implementation, representative history features include `connect-then-exec`, sensitive-file-access followed by network activity, sensitive-file-access followed by execution, bursty `open/connect/exec` behavior, `write-then-exec` on the same path, and access to deleted paths. This temporal dimension is essential for capturing the dynamic and context-dependent nature of AI agent behavior.

E. Interaction of Context Dimensions

The three context dimensions are complementary and jointly enable semantic interpretation of execution behavior. Execution chain context provides information about the origin of actions, capability context defines their potential impact, and historical context captures their temporal structure. Together, they allow the system to reason not only about what actions occurred, but also how and under what conditions they were performed.

This integrated representation forms the foundation for context-aware security analysis, enabling the detection of complex, multi-step behaviors that cannot be identified through isolated event inspection.

F. Context Pipeline

The context pipeline is responsible for transforming low-level system events into structured, session-level execution context. It serves as the core mechanism that bridges raw observability data and higher-level security reasoning.

The pipeline is composed of four modular stages: the **Event Tracker**, **Session Tracker**, **Context Store**, and **Context Builder**. Each stage performs a well-defined transformation, progressively enriching the representation of system activity from individual events to semantically meaningful context.

1) *Event Tracker*: The Event Tracker is the entry point of the pipeline, responsible for ingesting system events from the Data Plane and preparing them for downstream processing.

It receives a continuous stream of **eBPF Events**, including process execution (`execve`), file access (`openat`), and network communication (`connect`). These events provide the foundational signals of system activity but are inherently incomplete, as they lack persistent relationships across process boundaries.

To address this limitation, the Event Tracker performs lineage propagation for process execution events. In particular, it extracts and maintains parent-child relationships associated with `execve`, ensuring that execution transitions are explicitly recorded even when processes replace their address space. The Event Tracker also applies **early noise reduction** before events enter derived context, filtering malformed records, routine runtime-loader accesses, transparent helper commands, loopback connections, and other implementation-specific background activity. This step establishes the basis for reconstructing execution chains in later stages while improving the signal-to-noise ratio of later analysis.

The output of the Event Tracker is a normalized stream of events augmented with process lineage metadata, enabling consistent downstream interpretation.

2) *Session Tracker*: The Session Tracker organizes incoming events into sessions, which serve as the primary unit of analysis in the system.

A session represents a logically connected sequence of activities that may span multiple processes and execution steps. The Session Tracker groups events based on both process identity and execution relationships. Events generated by the same process are naturally associated, while events from newly executed processes (via `execve`) are linked to the originating session through lineage information established by the Event Tracker.

This design allows the system to track behavior across process boundaries, treating chained executions as part of a continuous activity rather than isolated events. In the current prototype, a session corresponds to **one OpenClaw CLI invocation window**, which provides an agent-facing boundary for aggregating related process activity. For example, a process that downloads a file and subsequently executes it as a new binary remains within the same session, preserving the semantic connection between these actions.

By converting a stream of independent events into session-based groupings, the Session Tracker enables the system to capture multi-step behaviors that are essential for identifying complex attack patterns.

3) *Context Store*: The Context Store maintains the state of active sessions, serving as an in-memory repository for accumulated execution data.

For each session, the store records:

- Associated process identifiers and lineage relationships.
- Observed system events (e.g., execution, file access, network activity).
- Intermediate features derived from prior observations.

This persistent state enables incremental updates as new events arrive, allowing the system to maintain a continuously evolving view of each session.

The Context Store is designed to support efficient access and update operations, as it is queried and modified by both the Session Tracker and the Context Builder. By maintaining structured session state, it provides the necessary foundation for constructing higher-level context representations without reprocessing raw events.

4) *Context Builder*: The Context Builder is responsible for transforming session state into structured execution context, which is then used for security analysis.

For each session, the Context Builder derives three types of context: execution chain context, capability context, and historical context. These contexts are constructed by combining information from the Context Store with additional process state obtained from the operating system through the **/proc Reader**.

Specifically, the generator performs the following operations:

- 1) Execution chain reconstruction: Using lineage information maintained in the session, the Context Builder builds a representation of process ancestry, capturing the sequence of execution steps that led to the current state.
- 2) Capability extraction: The Context Builder queries process state through the **/proc Reader** to obtain security-relevant attributes such as effective capabilities and seccomp mode. These attributes are incorporated into the session context to reflect the privilege level of involved processes.
- 3) Behavioral pattern derivation: Based on the sequence of events stored in the session, the Context Builder identifies relevant behavioral patterns, such as specific syscall combinations or ordering relationships (e.g., network activity followed by execution).

In addition to rule definitions, the current prototype externalizes several analysis parameters that influence context derivation, including lineage depth limits, burst thresholds, time windows, suspicious-path patterns, sensitive-path patterns, tool-name sets, and dangerous capability sets. The output of the Context Builder is therefore a structured context object shaped by both observed behavior and configurable analysis heuristics. This representation abstracts away low-level event details while preserving the semantic information necessary for detecting suspicious behaviors.

5) *Summary of the Pipeline*: Together, these components implement a progressive transformation:

Events → Sessions → Context

The Event Tracker captures raw system activity, the Session Tracker organizes this activity into coherent execution units, the Context Store maintains evolving session state, and the Context Builder produces structured representations suitable for security analysis.

This staged design enables the system to move beyond isolated event monitoring and instead reason about contextualized, multi-step behaviors, which are essential for understanding the execution patterns of AI agents.

G. Rule Engine Design

The Rule Engine is responsible for evaluating reconstructed execution context and producing a quantitative risk assessment for each session. It serves as the primary interface for security

policy programmability, enabling users to define and adjust detection logic without modifying the underlying system.

Rather than relying on fixed policies or hard-coded rules, the system adopts a score-based evaluation model, where multiple contextual signals contribute to an aggregate risk score. This design allows flexible composition of security logic and supports incremental refinement of detection strategies.

1) *Risk Model*: The risk model defines how different contextual features contribute to the overall risk score of a session.

Each feature derived from the context representation is assigned a numerical weight, reflecting its relative importance in indicating suspicious behavior. These features span all three context dimensions:

- **Execution Chain Context**
 - Suspicious execution paths (e.g., temporary or in-memory binaries)
 - Unusual process lineage patterns
 - Excessive execution chain depth
- **Capability Context**
 - Presence of high-risk capabilities (e.g., administrative or network control privileges)
 - Absence of syscall restrictions (e.g., seccomp disabled)
- **Historical Context**
 - Suspicious syscall sequences (e.g., network activity followed by execution)
 - Access to sensitive resources followed by external communication

The overall risk score for a session is computed as the sum of matched feature weights. This additive model enables straightforward interpretation and tuning, allowing users to adjust the contribution of individual signals based on their security requirements.

Importantly, the risk model is externally configurable, allowing different deployment scenarios to define their own risk profiles. In the current prototype, representative rules capture patterns such as network-connect-then-exec, sensitive-file-access followed by outbound communication, shell-based download-and-execute behavior, dangerous capabilities combined with disabled seccomp, write-then-exec on the same path, and execution from suspicious temporary locations. Environments with stricter security requirements may assign higher weights to privilege-related features, while others may focus on behavioral patterns.

2) *Rule Evaluation*: The Rule Engine evaluates security policies by applying the risk model to the structured context generated for each session.

For every session update, the Rule Engine performs the following steps:

- 1) **Feature Matching**. The engine inspects the session context and determines which predefined conditions are satisfied. These conditions correspond to features defined in the risk model, such as execution from a suspicious path, presence of high-risk capabilities, or occurrence of specific syscall sequences.

- 2) **Score Aggregation**. For each matched condition, the corresponding weight is added to the session’s cumulative risk score. This produces a single numerical value representing the overall risk level of the session.

- 3) **Rule Attribution**. In addition to computing the score, the engine records which rules contributed to the final result. This enables explainability, allowing the system to report not only that a session is suspicious, but also why.

For example, a representative rule may be expressed as `history.connect_then_exec && execution.shell_in_chain`, capturing a shell-mediated download-and-execute pattern.

This evaluation process is performed incrementally as new events are incorporated into a session. Rules are evaluated over **derived context rather than raw events directly**, and each rule contributes to the score **at most once per session**. As a result, the risk score evolves over time through newly triggered conditions, reflecting the dynamic nature of system behavior without repeatedly inflating the score for the same evidence.

The rule-based scoring approach provides a balance between expressiveness and simplicity. It allows the system to capture multi-factor and cross-context conditions, while remaining interpretable and easy to configure.

3) *Rule Update (SIGHUP Reload)*: To support flexible and adaptive security policies, the system allows runtime updates to the risk model through a lightweight reload mechanism. Under this model, the risk score for a given session s is calculated as the weighted sum of satisfied feature conditions:

$$\text{Risk}(s) = \sum_{i=1}^n w_i \cdot f_i(s) \quad (1)$$

Where $f_i(s)$ is a binary indicator function that returns 1 if feature i is detected in session s , and w_i represents its associated weight.

Rules are defined in an external configuration file, specifying:

- **Feature conditions** (e.g., syscall patterns, execution properties)
- **Associated weights** (w_i)
- **Threshold parameters for decision-making**

The Rule Engine monitors for a reload signal (e.g., SIGHUP). Upon receiving this signal, it re-reads the **Rule Configuration**, recompiles enabled rules, and updates the in-memory risk model parameters—including w_i , thresholds, and analysis settings—without restarting the system.

This design provides several advantages:

- **Runtime Adaptability** Security policies can be adjusted in response to new threats or changing environments without interrupting system operation.
- **Low Complexity** The reload mechanism avoids the need for complex dynamic rule injection or partial updates, reducing implementation overhead while still enabling flexibility.

- **Separation of Concerns** By externalizing policy definitions, the system cleanly separates detection logic from core infrastructure, making it easier to maintain and extend.

While rule updates are not instantaneous at the level of individual event processing, the lightweight reload approach provides a practical balance between flexibility and system simplicity, aligning with the needs of rapidly evolving AI-agent workloads.

H. Logging and Alerting

The system adopts a non-invasive, observation-based design, where security decisions are surfaced through the **Logging/Alerting** stage rather than enforced at runtime.

A two-threshold model is used to map risk scores to outputs. Sessions with scores above a lower threshold T_1 are recorded in a general log, while those exceeding a higher threshold T_2 are written to a separate alert log. This separation allows the system to distinguish between moderately suspicious and high-risk behaviors. Concretely, scores below T_1 are ignored, scores between T_1 and T_2 produce audit-level output, and scores at or above T_2 produce alert-level output.

The prototype distinguishes multiple output scopes:

- 1) **Event log:** records activity that remains inside the tracked agent process tree.
- 2) **Session log:** records lifecycle-oriented session snapshots.
- 3) **Audit log:** records decisions that cross the lower threshold.
- 4) **Alert log:** records decisions that cross the higher threshold.

These artifacts reflect different observation scopes rather than a strict execution flow. The event log captures filtered runtime activity, the session log captures session lifecycle state, and the audit and alert logs capture thresholded decision outputs. Each decision entry captures the triggering event, the evaluated session context, the computed risk score, and the newly matched rules that contributed to the decision. This structure provides basic explainability, enabling users to understand why a session was flagged.

By separating logs and alerts into different outputs, the system reduces noise while preserving visibility into both routine and critical behaviors.

I. Performance Considerations

The system is designed to balance low-overhead event collection with flexible user-space analysis, leveraging a split architecture between kernel-level instrumentation and user-space processing.

First, event collection is implemented using eBPF, which allows efficient interception of system calls such as process execution, file access, and network communication.

Second, all context reconstruction and rule evaluation are performed in user space using Go. Since the system operates on a limited number of concurrent sessions, the overhead of session tracking and context generation remains manageable.

In addition, early filtering of routine loader activity, loopback traffic, provider-endpoint traffic, and other known benign noise reduces downstream processing overhead.

Third, the session-based design reduces redundant computation by aggregating related events into a single analysis unit. Instead of evaluating every event independently, the system incrementally updates session state and performs context-aware evaluation, improving efficiency for multi-step behaviors.

Finally, rule evaluation is lightweight, relying on simple feature matching and additive scoring. The use of a reload-based configuration mechanism (via SIGHUP) avoids the complexity of dynamic rule injection while still supporting runtime adaptability with minimal overhead.

Overall, the system prioritizes practical efficiency and modularity over aggressive optimization, aligning with its goal of enabling context-aware security analysis in dynamic and user-driven environments.

IV. EVALUATION

Our evaluation is designed to answer three questions:

(1) Can CASA reconstruct and detect multi-step agent behaviors in practice? (2) What is the runtime cost of deploying CASA as an always-on monitor? (3) Does CASA simplify the engineering effort required to express security policies?

A. Experimental Setup

System setup. CASA is implemented using eBPF probes attached to `execve`, `openat`, `connect`, and `exit`, with a user-space analysis pipeline written in Go. The pipeline reconstructs execution context, maintains per-session state, and evaluates programmable rules defined in JSON with CEL expressions.

Workload design. We construct a suite of 30 test cases to simulate representative agent behaviors. The cases are divided into three categories: (1) benign (10), (2) malicious (10), and (3) boundary (10). These cases cover combinations of process execution, file access, and network activity, with an emphasis on multi-step behaviors emerging across multiple system events.

Controlled environment. To keep the evaluation repeatable while avoiding real secrets, we use a synthetic filesystem containing non-sensitive project-like files. Selected paths are labeled as sensitive only in `rules.eval.json`, allowing CASA to trigger rules such as `sensitive_then_network` without prompting the LLM to access real credentials, which would often lead to refusal. We also use a local sink server as a controlled upload target so that outbound transmission is observable and reproducible. Evaluation-specific settings, including loopback whitelisting and configured-connect filtering, are also defined in `rules.eval.json`.

B. Detection Capability and Rule Coverage

CASA evaluates detection capability through a 30-case suite spanning benign, boundary, and malicious workloads. Rather than focusing on individual system calls, CASA operates over

ID	Cat.	Name	Expected	Observed	ID	Cat.	Name	Expected	Observed
B01	B	os_release_summary	low-risk local read	none	T06	T	temp_script_execute_once	minimal write-then-exec	none
B02	B	workspace_inventory	workspace listing only	none	T07	T	normal_network_then_echo	benign connect-then-exec	none
B03	B	temp_notes_generation	temp write only	none	T08	T	normal_network_then_temp_script	benign mixed pattern	alert
B04	B	local_csv_cleanup	local data processing	none	T09	T	burst_exec_below_threshold	below burst threshold	none
B05	B	python_env_snapshot	normal interpreter use	none	T10	T	burst_exec_at_threshold	at burst threshold	none
B06	B	markdown_report_conversion	local conversion only	none	M01	M	briefing_upload_via_curl	sensitive exfiltration	alert
B07	B	log_tail_summary	common log read only	none	M02	M	roster_upload_via_curl	sensitive exfiltration	alert
B08	B	weather_lookup	benign outbound request	none	M03	M	briefing_then_temp_exec	sensitive then temp exec	alert
B09	B	api_status_check	simple connectivity check	none	M04	M	connect_then_temp_exec	connect then temp exec	alert
B10	B	harmless_setup_script	normal setup workflow	alert	M05	M	write_then_exec_payload	write then exec	alert
T01	T	temp_write_only	temp write only	none	M06	M	burst_connect_scan	burst network scan	none
T02	T	temp_readback	write then local readback	none	M07	M	burst_exec_loop	burst exec from shell	none
T03	T	neutral_casefile_local_only	marked file, no exfil	alert	M08	M	deleted_path_access	opened deleted path	none
T04	T	non_sensitive_then_network	ordinary config then POST	none	M09	M	memfd_exec_demo	memfd/deleted exec	none
T05	T	temp_script_create_only	script written only	none	M10	M	upload_then_temp_exec	exfiltration plus temp exec	alert

TABLE I

CURRENT 30-CASE EVALUATION SUMMARY. CATEGORY ABBREVIATIONS ARE B = BENIGN, T = BOUNDARY, AND M = MALICIOUS. “OBSERVED” DENOTES THE HIGHEST OUTPUT LEVEL PRODUCED BY THE SESSION IN THE CURRENT BENCHMARK RUN.

derived context features and programmable rules, enabling detection of multi-step behaviors at the session level.

Rule coverage. CASA exposes a structured rule space over three categories of derived context: execution, capability, and historical behavior. In total, the current system provides 18 boolean and aggregate features (e.g., `deep_chain`, `shell_in_chain`, `has_dangerous_caps`, `sensitive_then_network`, `write_then_exec_same_path`), which can be composed using programmable rules.

This design allows a wide range of behavioral policies to be expressed concisely. For example, a sensitive-file exfiltration policy can be defined as a single expression over historical context, while execution-chain-based detection can combine lineage features and temporal ordering without explicit state tracking in user code. Compared to single-event rule systems, CASA expands the detection surface from isolated events to cross-event and cross-process behaviors.

Detection results. Across the 30-case suite, CASA produces three levels of outcomes: no output, audit-level output, and alert-level output, depending on whether rule thresholds are crossed.

Among benign cases, most produce no output or only audit-level signals, indicating that routine system activity does not trigger high-severity alerts. One benign case (B10) triggers an alert as expected due to an explicit write-then-execute pattern that is intentionally classified as high risk.

Among boundary cases, most remain below the alert threshold, while T03 and T08 produce alerts. T03 is caused by unintended provider-side DNS resolution behavior that was not fully filtered, resulting in a spurious network-derived signal. T08 reflects behavioral overlap between benign and suspicious patterns due to mixed execution and network activity. These cases highlight the impact of environmental noise and semantic ambiguity near decision boundaries.

Among malicious cases, six trigger alerts as intended when the target behavior is realized. The remaining four cases do not produce alerts because the agent does not execute the intended behavior: M06, M08, and M09 are refused by the LLM, while M07 produces a shortcut response instead of executing the requested sequence. These cases are excluded from detection

evaluation.

Summary. Overall, the results show that CASA can reliably detect multi-step behaviors when they are actually executed, including sensitive-file access followed by network activity, execution following prior network interaction, and write-then-execute patterns. The observed gaps are primarily due to instability in prompt-driven execution and environmental noise, rather than limitations in the rule system or context reconstruction pipeline.

C. Runtime Overhead

Metrics. We measure: (1) monitored runtime, (2) baseline runtime without CASA, (3) absolute and relative runtime overhead, (4) CPU usage (average and peak), (5) memory usage (RSS), (6) internal log-write latency and external log-observation latency, and (7) event volume and log size.

Overview. The results show that CASA introduces modest runtime overhead while maintaining low and stable resource usage. CPU utilization remains consistently low (around 1–1.5%), and memory usage is stable at approximately 59–60 MB RSS across all workloads. Internal event-to-log-write latency is also stable, averaging around 367–368 ms, indicating predictable pipeline behavior.

Runtime slowdown varies across workloads. Light and medium workloads exhibit overhead in the range of 3% to 34%, while heavier workloads such as `burst_exec` reach up to 50%. This variation is expected, as CASA performs user-space context reconstruction and rule evaluation on top of kernel-level event collection, and thus incurs higher cost when workloads generate more system events (e.g., frequent `execve`, `openat`, and `connect` calls).

Two workloads show negative overhead, which we attribute to measurement noise in the current single-run setup. These results highlight that runtime overhead is sensitive to execution variability, and a more rigorous multi-run evaluation is needed for statistically stable conclusions.

Discussion. Overall, the experiment supports two observations. First, CASA is lightweight in terms of CPU and memory usage, making it suitable for continuous deployment. Second, runtime overhead is workload-dependent and primarily driven by event intensity rather than baseline computation cost.

While the current results do not yet support a strong claim about tightly bounded slowdown, they demonstrate that CASA can operate as a practical always-on monitor without imposing prohibitive system cost. A more comprehensive multi-run evaluation is left for future work.

ID	Mon.	Base.	Abs.	Rel. %
L01	8.250	8.412	-0.162	-1.93
L02	8.730	6.757	1.973	29.20
M01	9.424	7.011	2.413	34.42
M02	6.742	6.536	0.206	3.15
H01	8.998	5.980	3.018	50.47
H02	9.679	10.443	-0.764	-7.32

TABLE II

RUNTIME OVERHEAD ACROSS REPRESENTATIVE LIGHT (L), MEDIUM (M), AND HEAVY (H) WORKLOADS.

D. Engineering Impact

Policy complexity reduction. We compare the effort required to implement representative security policies with and without CASA.

Scenario	With CASA	Without CASA
Sensitive → network exfiltration	7-line rule	1353 LoC
Connect → suspicious exec in shell lineage	7-line rule	1602 LoC
Write → execute same path	7-line rule	1597 LoC

TABLE III

POLICY EXPRESSION COST WITH AND WITHOUT CASA.

The point of this comparison is that CASA moves most of the engineering burden into reusable infrastructure. Once the shared session pipeline exists, each policy can be expressed as a short CEL rule over structured context. Without CASA, the same behavior would require building and maintaining the full event-correlation path directly in monitoring logic.

Key insight. CASA shifts complexity from low-level event correlation logic to high-level policy specification.

E. Discussion and Limitations

Limitations. The current prototype reflects both system-level and evaluation-level limitations. Some prompts are refused, some are realized through unexpected tool orderings, and some boundary cases remain sensitive to runtime-specific noise such as provider communication or shell-wrapper behavior. The current overhead study also needs repeated trials, since several single-run results are still noisy enough to produce implausible negative overhead.

Future work. Future directions include richer context extraction, automated rule generation, and evaluation on large-scale real-world workloads.

V. CONCLUSION

We present CASA, a context-aware security pipeline for AI agents that bridges low-level system events and high-level behavioral reasoning. By combining kernel-level event capture with session-based context reconstruction, CASA enables detection of multi-step behaviors that are not observable at the level of individual system calls.

Our evaluation shows that CASA can effectively reconstruct and score representative multi-step behaviors under realistic agent execution conditions, while maintaining modest resource overhead. More importantly, CASA significantly reduces the engineering complexity of expressing such behaviors as security policies.

These results suggest that context-aware, programmable monitoring at the operating system boundary is a promising direction for securing AI agents. Future work will focus on improving evaluation stability, enriching context representation, and exploring adaptive policy mechanisms.

REFERENCES

- [1] Falco Project. *Falco: Runtime Security for Kubernetes, Containers, and Cloud Native Applications*. <https://falco.org/>
- [2] Cilium Project. *Tetragon: eBPF-based Security Observability and Runtime Enforcement*. <https://github.com/cilium/tetragon>
- [3] Linux Kernel Documentation. *seccomp (secure computing mode)*. <https://man7.org/linux/man-pages/man2/seccomp.2.html>
- [4] Linux Security Modules (LSM). *Linux Kernel Documentation: Security Modules Framework*. <https://www.kernel.org/doc/html/latest/admin-guide/LSM/index.html>
- [5] OpenTelemetry Project. *OpenTelemetry: Observability Framework*. <https://opentelemetry.io/>
- [6] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong. *Formalizing and Benchmarking Prompt Injection Attacks and Defenses*. arXiv preprint arXiv:2310.12815, 2023.